

stats_in_julia_solutions

April 4, 2017

1 Question 1

Using the data in `chronic_kidney_disease.csv`, determine whether or not the oldest patient in the sample has chronic kidney disease. Note that the `class` variable indicates a patient's CKD status.

As a hint, you will probably want to use the `maximum()` function and the `find()` function.

And you'll need to consider how to handle NA values; there is a function `dropna()` that will be useful.

```
In [37]: using DataFrames
```

```
ckd = readtable("../data/chronic_kidney_disease.csv");
```

```
In [10]: maxage = maximum(dropna(ckd[:age]))
```

```
oldest_idx = find(ckd[:age] .== maxage)
```

```
ckd[oldest_idx, :class]
```

```
Out[10]: 1-element DataArrays.DataArray{String,1}:  
         "ckd"
```

2 Question 2

Using the `chronic_kidney_disease.csv` dataset from above, use an appropriate statistical test to determine if patients with chronic kidney disease (CKD) have significantly higher blood urea than patients without CKD.

As a hint, you will likely want to use the `HypothesisTests` package.

```
In [10]: using DataFrames  
         using HypothesisTests
```

```
ckd = readtable("../data/chronic_kidney_disease.csv")
```

```
# get row indices
```

```
ckd_ids = find(ckd[:class] .== "ckd")
```

```
nonckd_ids = find(ckd[:class] .== "notckd")
```

```

# get `blood_urea` scores
ckd_grp = ckd[ckd_ids, :blood_urea]
nonckd_grp = ckd[nonckd_ids, :blood_urea]

UnequalVarianceTTest(dropna(ckd_grp), dropna(nonckd_grp))

Out[10]: Two sample t-test (unequal variance)
-----
Population details:
  parameter of interest:   Mean difference
  value under h_0:         0
  point estimate:          39.590418424753864
  95% confidence interval: (31.86529602463314, 47.31554082487459)

Test summary:
  outcome with 95% confidence: reject h_0
  two-sided p-value:          1.7990230669795845e-20 (extremely significant)

Details:
  number of observations:   [237,144]
  t-statistic:              10.090709649973618
  degrees of freedom:       264.8784441340254
  empirical standard error: 3.9234523435977935

```

3 Question 3

Using the `chronic_kidney_disease.csv` data, determine which of the following predictors (if any) are related chronic kidney disease (CKD):

```

blood_urea,
hemoglobin,
red_blood_cell_count,
white_blood_cell_count.

```

Hints: - There are a few ways this could be done, let's use a single regression model of some kind - Our outcome variable is `class` in the CKD data, this needs to be re-coded as 0/1

```

In [6]: using GLM
        using DataFrames

ckd = readtable("../data/chronic_kidney_disease.csv", makefactors = true)

ckd[:has_ckd] = [x == "ckd" ? 1 : 0 for x in ckd[:class]]

mod3 = @formula(has_ckd ~ 1 + blood_urea + hemoglobin + red_blood_cell_count)

fm3 = glm(mod3, ckd, Binomial())

```

```
Out [6]: DataFrames.DataFrameRegressionModel{GLM.GeneralizedLinearModel{GLM.GlmResp}

Formula: has_ckd ~ 1 + blood_urea + hemoglobin + red_blood_cell_count + whi

Coefficients:

                Estimate   Std. Error   z value   Pr(>|z|)
(Intercept)         20.5118       3.82243    5.36617    <1e-7
blood_urea           0.0135445    0.0152221    0.889787    0.3736
hemoglobin          -1.4337       0.273702   -5.23817    <1e-6
red_blood_cell_count -0.685826    0.399729   -1.71573    0.0862
white_blood_cell_count 0.000152042 0.000141559  1.07406    0.2828
```

4 Question 4

Using the `stagec` data from above, fit several models experimenting with different numbers of trees and different numbers of variable subsets for candidate splitting (i.e., `m_try`, the third argument to the `build_forest()` function).

In order to evaluate the quality of the models on training data, write a function that calculates the mean-squared error of the fitted model. The function should take 3 arguments: (1) the fitted model, (2) the vector with the outcome variable and (3) the matrix of predictors.

What was the mean-squared error of your best-fitting model?

```
In [11]: # This is a quick function to obtain the mean-squared
# error of a fitted random forest (or bagged tree) model.
```

```
function mse(fitted, y, X)
    yhat = apply_forest(fitted, X)
    sqerr = (y .- yhat).^2
    out = mean(sqerr)
    return out
end
```

```
Out[11]: mse (generic function with 1 method)
```

```
In [14]: # Load the data
using DecisionTree
using RDatasets
stagec = dataset("rpart", "stagec")

stagec_comp = stagec[completecases(stagec), :];
```

```
WARNING: using DecisionTree.fit! in module Main conflicts with an existing identifier
WARNING: using DecisionTree.R2 in module Main conflicts with an existing identifier
WARNING: using DecisionTree.predict in module Main conflicts with an existing identifier
```

```
In [15]: # Clean up data and fit model
```

```

is_tetraploid = stagec_comp[:Ploidy] .== "tetraploid"

stagec_comp[:tetra] = is_tetraploid

# must convert to Array
y = convert(Array{Float64,1}, stagec_comp[:G2])
X = convert(Array{Float64,2}, stagec_comp[:, :Age, :Grade, :Gleason, :EET, :

fm4a = build_forest(y, X, 5, 100)
fm4b = build_forest(y, X, 3, 100)
fm4c = build_forest(y, X, 5, 500)

```

```

Out[15]: Ensemble of Decision Trees
Trees:      500
Avg Leaves: 31.272
Avg Depth:  9.914

```

```

In [16]: @show mse(fm4a, y, X)
          @show mse(fm4b, y, X)
          @show mse(fm4c, y, X)

```

```

mse(fm4a, y, X) = 18.4084866706618
mse(fm4b, y, X) = 22.123257454931814
mse(fm4c, y, X) = 18.49337910973606

```

```

Out[16]: 18.49337910973606

```

5 Question 5

Using the `aldh2` dataset from the `gap` package in R, try fitting a few random forest (or bagged tree) models using Julia to predict whether a given patient is an alcoholic using their genetic information.

What is the prediction accuracy of your best model? What were the meta-parameters of your best-fitting model?

The data can be loaded using the code below.

```

In [30]: using RDatasets
          using DecisionTree

          aldh2 = dataset("gap", "aldh2");

In [33]: # must convert to Array
          y = convert(Array{Float64,1}, aldh2[:,Y])
          X = convert(Array{Float64,2}, aldh2[:, 3:end])

          # fit model
          fm10 = build_forest(y, X, 10, 500)

```

```
Out[33]: Ensemble of Decision Trees
Trees:      500
Avg Leaves: 64.718
Avg Depth:  14.248
```

```
In [34]: # get prediction accuracy
yhat_bool = apply_forest(fm10, X) .> 0.5
ybool = y .== 1
mean(ybool .== yhat_bool)
```

```
Out[34]: 0.9809885931558935
```

6 Question 6

Use R and the `randomForest` package via the `RCall.jl` package in Julia to fit a random forest model on the `chronic_kidney_disease.csv` data set. In particular, fit a model with 5000 trees to predict whether or not patients have chronic kidney disease.

After fitting the model, extract the variable importance estimates (mean Gini decrease) from the fitted model. Pass a dataframe back to Julia that has two columns (1) name of the predictor, and (2) mean Gini decrease for that predictor.

Sort the returned data frame such that larger values are at the top.

The following steps should serve as a general to complete this:

1. Read the data in to Julia
2. Ensure `RCall.jl` is loaded
3. Pass the dataframe from Julia to R
4. Fit the model in R using `randomForest()` function with argument `ntrees = 5000`
5. Use the `importance()` function to extract the estimates of variable importance from the fitted model
6. Pass the estimates back to Julia

Some Hints:

- The `importance()` function in R returns a data frame whose row names are the variable names, and the last column is mean Gini decrease
- The documentation for the `randomForest` package in R can be found here: <https://cran.r-project.org/web/packages/randomForest/randomForest.pdf>
- The `sortperm()` function in Julia will be useful for sorting in the last step

```
In [35]: using DataFrames
using RCall

ckd = readtable("../data/chronic_kidney_disease.csv", makefactors = true)
```

```

@rput ckd

R"
library(randomForest)

ckd2 <- ckd[complete.cases(ckd), ]

fm6 <- randomForest(class ~ ., ckd2, importance = TRUE, ntree = 5000)

imp_df <- importance(fm6)

gini_decrease <- imp_df[, ncol(imp_df)]

preds <- row.names(imp_df)

var_imp_df <- data.frame(preds, gini_decrease)
"
@rget var_imp_df;
In [36]: indcs = sortperm(var_imp_df[:, 2], rev = true)
res = var_imp_df[indcs, :]
Out [36]: 24x2 DataFrames.DataFrame

```

Row	preds	gini_decrease
1	"albumin"	10.5288
2	"serum_creatinine"	8.49594
3	"packed_cell_volume"	8.46403
4	"hemoglobin"	8.42965
5	"red_blood_cell_count"	6.2716
6	"specific_gravity"	5.00724
7	"blood_urea"	3.82916
8	"hypertension"	3.46274
9	"pus_cell"	1.66581
10	"diabetes_mellitus"	1.44764
11	"blood_glucose_random"	1.12767
12	"sodium"	0.8763
13	"white_blood_cell_count"	0.624537
14	"sugar"	0.390784
15	"pedal_edema"	0.339301
16	"red_blood_cells"	0.298891
17	"blood_pressure"	0.290002
18	"appetite"	0.238108
19	"anemia"	0.148841
20	"pus_cell_clumps"	0.0765711
21	"bacteria"	0.0746202
22	"potassium"	0.0720374
23	"age"	0.0613
24	"coronary_artery_disease"	0.0250922

```
In [ ]:
```